

```

-----import libraries-----
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import KFold,cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report,accuracy_score
from sklearn.preprocessing import StandardScaler,RobustScaler
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix,classification_report

```

```

-----read dataset-----
data1 = pd.read_csv('Downloads/SalaryData_Train(1).csv')
data1.head(3)

```

|   | age | workclass        | education | educationno | maritalstatus      | occupation        | relationship  | race  | sex  | capitalgain | capitalloss | hoursperweek | native        | Salary |
|---|-----|------------------|-----------|-------------|--------------------|-------------------|---------------|-------|------|-------------|-------------|--------------|---------------|--------|
| 0 | 39  | State-gov        | Bachelors | 13          | Never-married      | Adm-clerical      | Not-in-family | White | Male | 2174        | 0           | 40           | United-States | <=50K  |
| 1 | 50  | Self-emp-not-inc | Bachelors | 13          | Married-civ-spouse | Exec-managerial   | Husband       | White | Male | 0           | 0           | 13           | United-States | <=50K  |
| 2 | 38  | Private          | HS-grad   | 9           | Divorced           | Handlers-cleaners | Not-in-family | White | Male | 0           | 0           | 40           | United-States | <=50K  |

```

data2 = pd.read_csv('Downloads/SalaryData_Test(1).csv')
data2.head(3)

```

|   | age | workclass | education  | educationno | maritalstatus      | occupation        | relationship | race  | sex  | capitalgain | capitalloss | hoursperweek | native        | Salary |
|---|-----|-----------|------------|-------------|--------------------|-------------------|--------------|-------|------|-------------|-------------|--------------|---------------|--------|
| 0 | 25  | Private   | 11th       | 7           | Never-married      | Machine-op-inspct | Own-child    | Black | Male | 0           | 0           | 40           | United-States | <=50K  |
| 1 | 38  | Private   | HS-grad    | 9           | Married-civ-spouse | Farming-fishing   | Husband      | White | Male | 0           | 0           | 50           | United-States | <=50K  |
| 2 | 28  | Local-gov | Assoc-acdm | 12          | Married-civ-spouse | Protective-serv   | Husband      | White | Male | 0           | 0           | 40           | United-States | >50K   |

```

-----shape-----

```

```
data1.shape
```

```
(30161, 14)
```

```
data2.shape
```

```
(15060, 14)
```

```

-----info-----

```

```
data1.info()
```

```

RangeIndex: 30161 entries, 0 to 30160
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  ---
0    age                   30161 non-null  int64
1    workclass             30161 non-null  object
2    education             30161 non-null  object
3    educationno           30161 non-null  int64
4    maritalstatus         30161 non-null  object
5    occupation            30161 non-null  object
6    relationship          30161 non-null  object
7    race                  30161 non-null  object
8    sex                   30161 non-null  object
9    capitalgain           30161 non-null  int64
10   capitalloss           30161 non-null  int64
11   hoursperweek          30161 non-null  int64
12   native                30161 non-null  object
13   Salary                30161 non-null  object
dtypes: int64(5), object(9)
memory usage: 3.2+ MB

```

-----describe-----  
data1.describe()

|       | age          | educationno  | capitalgain  | capitalloss  | hoursperweek |
|-------|--------------|--------------|--------------|--------------|--------------|
| count | 30161.000000 | 30161.000000 | 30161.000000 | 30161.000000 | 30161.000000 |
| mean  | 38.438115    | 10.121316    | 1092.044064  | 88.302311    | 40.931269    |
| std   | 13.134830    | 2.550037     | 7406.466611  | 404.121321   | 11.980182    |
| min   | 17.000000    | 1.000000     | 0.000000     | 0.000000     | 1.000000     |
| 25%   | 28.000000    | 9.000000     | 0.000000     | 0.000000     | 40.000000    |
| 50%   | 37.000000    | 10.000000    | 0.000000     | 0.000000     | 40.000000    |

data2.info()

```

#   Column                Non-Null Count  Dtype
---  ---
0    age                   15060 non-null  int64
1    workclass             15060 non-null  object
2    education             15060 non-null  object
3    educationno           15060 non-null  int64
4    maritalstatus         15060 non-null  object
5    occupation            15060 non-null  object
6    relationship          15060 non-null  object
7    race                  15060 non-null  object
8    sex                   15060 non-null  object
9    capitalgain           15060 non-null  int64
10   capitalloss           15060 non-null  int64
11   hoursperweek          15060 non-null  int64
12   native                15060 non-null  object
13   Salary                15060 non-null  object
dtypes: int64(5), object(9)
memory usage: 1.6+ MB

```

data2.describe()

|       | age          | educationno  | capitalgain  | capitalloss  | hoursperweek |
|-------|--------------|--------------|--------------|--------------|--------------|
| count | 15060.000000 | 15060.000000 | 15060.000000 | 15060.000000 | 15060.000000 |
| mean  | 38.768327    | 10.112749    | 1120.301594  | 89.041899    | 40.951594    |
| std   | 13.380676    | 2.558727     | 7703.181842  | 406.283245   | 12.062831    |
| min   | 17.000000    | 1.000000     | 0.000000     | 0.000000     | 1.000000     |
| 25%   | 28.000000    | 9.000000     | 0.000000     | 0.000000     | 40.000000    |
| 50%   | 37.000000    | 10.000000    | 0.000000     | 0.000000     | 40.000000    |
| 75%   | 48.000000    | 13.000000    | 0.000000     | 0.000000     | 45.000000    |

```
data1.workclass.unique()
array([' State-gov', ' Self-emp-not-inc', ' Private', ' Federal-gov',
       ' Local-gov', ' Self-emp-inc', ' Without-pay'], dtype=object)
```

```
data1.workclass.value_counts()

Private                22285
Self-emp-not-inc      2499
Local-gov              2067
State-gov              1279
Self-emp-inc           1074
Federal-gov            943
Without-pay            14
Name: workclass, dtype: int64
```

```
data1.occupation.unique()
array([' Adm-clerical', ' Exec-managerial', ' Handlers-cleaners',
       ' Prof-specialty', ' Other-service', ' Sales', ' Transport-moving',
       ' Farming-fishing', ' Machine-op-inspct', ' Tech-support',
       ' Craft-repair', ' Protective-serv', ' Armed-Forces',
       ' Priv-house-serv'], dtype=object)
```

```
data1.occupation.value_counts()

Prof-specialty        4038
Craft-repair          4030
Exec-managerial        3992
Adm-clerical           3721
Sales                  3584
Other-service          3212
Machine-op-inspct     1965
```

```
data1.native.unique()
array([' United-States', ' Cuba', ' Jamaica', ' India', ' Mexico',
       ' Puerto-Rico', ' Honduras', ' England', ' Canada', ' Germany',
       ' Iran', ' Philippines', ' Poland', ' Columbia', ' Cambodia',
       ' Thailand', ' Ecuador', ' Laos', ' Taiwan', ' Haiti', ' Portugal',
       ' Dominican-Republic', ' El-Salvador', ' France', ' Guatemala',
       ' Italy', ' China', ' South', ' Japan', ' Yugoslavia', ' Peru',
       ' Outlying-US(Guam-USVI-etc)', ' Scotland', ' Trinidad&Tobago',
       ' Greece', ' Nicaragua', ' Vietnam', ' Hong', ' Ireland',
       ' Hungary'], dtype=object)
```

```
data1.native.value_counts()
```

|               |       |
|---------------|-------|
| United-States | 27504 |
| Mexico        | 610   |
| Philippines   | 188   |
| Germany       | 128   |
| Puerto-Rico   | 109   |
| Canada        | 107   |
| El-Salvador   | 100   |
| India         | 100   |
| Cuba          | 92    |
| England       | 86    |
| Jamaica       | 80    |
| South         | 71    |
| Italy         | 68    |
| China         | 68    |

```
numerical = [var for var in data1.columns if data1[var].dtypes!='0']
print('numerical values are',numerical)
```

```
x = data1.drop(['Salary'],axis=1)
y = data1['Salary']
```

-----traintest-----

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
x_train.shape,x_test.shape,y_train.shape,y_test.shape

((21112, 13), (9049, 13), (21112,), (9049,))
```

x\_train.dtypes

|               |        |
|---------------|--------|
| age           | int64  |
| workclass     | object |
| education     | object |
| educationno   | int64  |
| maritalstatus | object |
| occupation    | object |
| relationship  | object |
| race          | object |
| sex           | object |
| capitalgain   | int64  |
| capitalloss   | int64  |
| hoursperweek  | int64  |
| native        | object |
| dtype:        | object |

x\_test.dtypes

|               |        |
|---------------|--------|
| age           | int64  |
| workclass     | object |
| education     | object |
| educationno   | int64  |
| maritalstatus | object |
| occupation    | object |
| relationship  | object |
| race          | object |
| sex           | object |
| capitalgain   | int64  |
| capitalloss   | int64  |

```
categorical = [col for col in x_train.columns if x_train[col].dtypes=='O']
categorical
```

```
['workclass',
 'education',
 'maritalstatus',
 'occupation',
 'relationship',
 'race',
 'sex',
 'native']
```

```
numerical = [col for col in x_train.columns if x_train[col].dtypes=='O']
numerical
```

---

```
['age', 'educationno', 'capitalgain', 'capitalloss', 'hoursperweek']
```

```
x_train[categorical].isnull().mean()
```

```
workclass      0.0
education      0.0
maritalstatus  0.0
occupation     0.0
relationship   0.0
race           0.0
sex            0.0
native         0.0
dtype: float64
```

```
for col in categorical:
```

```
    if x_train[col].isnull().mean()>0:
        print(col,(x_train[col].isnull().mean()))
```

```
for df2 in [x_train,x_test]:
```

```
    df2['workclass'].fillna(x_train['workclass'].model()[0],inplace=True)
    df2['occupation'].fillna(x_train['occupation'].model()[0],inplace=True)
    df2['native'].fillna(x_train['native'].model()[0],inplace=True)
```

```
x_train[categorical].isnull().sum()
```

```
workclass      0
education      0
maritalstatus  0
occupation     0
relationship   0
race           0
sex            0
native         0
dtype: int64
```

```
x_test[categorical].isnull().sum()
```

```
workclass      0
education      0
maritalstatus  0
occupation     0
relationship   0
race           0
sex            0
native         0
dtype: int64
```

---

```
x_train.isnull().sum()
```

```
age           0
workclass     0
education     0
educationno   0
maritalstatus 0
occupation   0
relationship  0
race          0
sex           0
capitalgain   0
capitalloss   0
hoursperweek  0
native        0
dtype: int64
```

---

```
x_test.isnull().sum()
```

```
age           0
workclass     0
education     0
educationno   0
maritalstatus 0
occupation   0
relationship  0
race          0
sex           0
capitalgain   0
capitalloss   0
hoursperweek  0
native        0
dtype: int64
```

```
pip install category_encoders
```

```
Collecting category_encoders
  Downloading category_encoders-2.3.0-py2.py3-none-any.whl (82 kB)
Requirement already satisfied: patsy>=0.5.1 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (0.5.1)
Requirement already satisfied: scikit-learn>=0.20.0 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (0.24.1)
Requirement already satisfied: statsmodels>=0.9.0 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (0.12.2)
Requirement already satisfied: scipy>=1.0.0 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (1.6.2)
Requirement already satisfied: pandas>=0.21.1 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (1.2.4)
Requirement already satisfied: numpy>=1.14.0 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (1.20.1)
Requirement already satisfied: pytz>=2017.3 in c:\users\dell\anaconda3\lib\site-packages (from pandas>=0.21.1->category_encoders) (2021.1)
Requirement already satisfied: python-dateutil>=2.7.3 in c:\users\dell\anaconda3\lib\site-packages (from pandas>=0.21.1->category_encoders) (2.8.1)
Requirement already satisfied: six in c:\users\dell\anaconda3\lib\site-packages (from patsy>=0.5.1->category_encoders) (1.15.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\dell\anaconda3\lib\site-packages (from scikit-learn>=0.20.0->category_encoders) (2.1.0)
Requirement already satisfied: joblib>=0.11 in c:\users\dell\anaconda3\lib\site-packages (from scikit-learn>=0.20.0->category_encoders) (1.0.1)
Installing collected packages: category-encoders
Successfully installed category-encoders-2.3.0
Note: you may need to restart the kernel to use updated packages.
```

```
import category_encoders as ce
encoder =
ce.OneHotEncoder(cols=['workclass','education','maritalstatus','occupation','relationship','race','sex','native'])
x_train = encoder.fit_transform(x_train)
x_test = encoder.fit_transform(x_test)
```

```
x_train.head()
```

|      | age | workclass_1 | workclass_2 | workclass_3 | workclass_4 | workclass_5 | workclass_6 | workclass_7 | education_1 | education_2 |
|------|-----|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| 8166 | 54  | 1           | 0           | 0           | 0           | 0           | 0           | 0           | 1           | 0           |
| 7138 | 21  | 0           | 1           | 0           | 0           | 0           | 0           | 0           | 1           | 0           |
| 437  | 30  | 0           | 1           | 0           | 0           | 0           | 0           | 0           | 0           | 1           |
| 5436 | 42  | 0           | 1           | 0           | 0           | 0           | 0           | 0           | 0           | 1           |
| 6541 | 37  | 0           | 0           | 1           | 0           | 0           | 0           | 0           | 0           | 1           |

```
rows x 102 columns
```

```
x_train.shape
```

```
(21112, 102)
```

```
x_test.head()
```

|       | age | workclass_1 | workclass_2 | workclass_3 | workclass_4 | workclass_5 | workclass_6 | workclass_7 |
|-------|-----|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| 25338 | 21  | 0           | 1           | 0           | 0           | 0           | 0           | 0           |
| 18840 | 21  | 0           | 1           | 0           | 0           | 0           | 0           | 0           |
| 8391  | 56  | 0           | 1           | 0           | 0           | 0           | 0           | 0           |
| 18258 | 43  | 1           | 0           | 0           | 0           | 0           | 0           | 0           |

```
cols = x_train.columns
```

```
scaler = RobustScaler()
```

```
x_train = scaler.fit_transform(x_train)
```

```
x_test = scaler.fit_transform(x_test)
```

```
x_train = pd.DataFrame(x_train,columns=[cols])
```

```
x_test = pd.DataFrame(x_test,columns=[cols])
```

```
x_train.head()
```

```
-----gaussian-----
```

```
gnb = GaussianNB()
```

```
gnb.fit(x_train,y_train)
```

```
GaussianNB()
```

```
ypredict = gnb.predict(x_test)
ypredict
array([' >50K', ' <=50K', ' <=50K', ..., ' >50K', ' <=50K', ' <=50K'],
      dtype='<U6')
```

```
print('Model accuracy score:{0:0.4f}'.format(accuracy_score(y_test,ypredict)))
```

```
Model accuracy score:{0:0.4f} 0.7515747596419494
```

```
ypredict_train = gnb.predict(x_train)
ypredict_train
```

```
array([' >50K', ' >50K', ' <=50K', ..., ' <=50K', ' <=50K', ' <=50K'],
      dtype='<U6')
```

```
print('Training - set accuracy score:{0:0.4f}'.format(accuracy_score(y_train,ypredict_train)))
```

```
Training - set accuracy score:{0:0.4f} 0.7975085259568018
```

```
y_test.value_counts()
```

```
<=50K    6819
>50K      2230
Name: Salary, dtype: int64
```

```
cm = confusion_matrix(y_test,ypredict)
print('Confusion Matrix \n\n',cm)
print('\n True Positives(TP)=',cm[0,0])
print('\n True Negatives(TN)=',cm[1,1])
print('\n False Positives(FP)=',cm[0,1])
print('\n False Negatives(FN)=',cm[1,0])
```

```
[[ 5113  1706]
 [  542  1688]]
```

```
True Positives(TP)= 5113
```


```
True Negatives(TN)= 1688
```

```
False Positives(FP)= 1706
```

```
False Negatives(FN)= 542
```



```
cm_matrix = pd.DataFrame(data=cm,columns=['actual positive 1','actual negative 0'],index=['predict positive 1','predict negative 0'])
sns.heatmap(cm_matrix,annot=True,fmt='d')
```



```
print(classification_report(y_test,ypredict))
```

|              |      |      |      |      |
|--------------|------|------|------|------|
| <=50K        | 0.90 | 0.75 | 0.82 | 6819 |
| >50K         | 0.50 | 0.76 | 0.60 | 2230 |
| accuracy     |      |      | 0.75 | 9049 |
| macro avg    | 0.70 | 0.75 | 0.71 | 9049 |
| weighted avg | 0.80 | 0.75 | 0.77 | 9049 |

```
TP = cm[0,0]
```

```
TN = cm[1,1]
```

```
FP = cm[0,1]
```

```
FN = cm[1,0]
```

```
classification_accuracy = (TP+TN)/float(TP+TN+FP+FN)
```

```
print('classification accuracy:{0:0.4f}'.format(classification_accuracy))
```

```
classification accuracy:{0:0.4f} 0.7515747596419494
```

```
classification_error = (FP+FN)/float(TP+TN+FP+FN)
```

```
print('classification error:{0:0.4f}'.format(classification_error))
```

---

```
classification error:{0:0.4f} 0.2484252403580506
```

```
precision = TP/float(TP+FP)
print('precision:{0:0.4f}',format(precision))
```

```
precision:{0:0.4f} 0.749816688664027
```

```
recall = TP/float(TP+FN)
print('recall:{0:0.4f}',format(recall))
```

```
recall:{0:0.4f} 0.9041556145004421
```

```
true_positive_rate = TP/float(TP+FN)
print('true_positive_rate:{0:0.4f}',format(true_positive_rate))
```

```
true_positive_rate:{0:0.4f} 0.9041556145004421
```

```
false_positive_rate = FP/float(FP+TN)
print('false_positive_rate:{0:0.4f}',format(false_positive_rate))
```

```
false_positive_rate:{0:0.4f} 0.502651738361815
```

```
specificity = TN/float(FP+TN)
print('specificity:{0:0.4f}',format(specificity))
```

```
specificity:{0:0.4f} 0.497348261638185
```

```
ypredict_prob = gnb.predict_proba(x_test)[0:10]
ypredict_prob
```


```
array([[5.95373694e-02, 9.40462631e-01],
       [9.99992594e-01, 7.40639066e-06],
       [9.99865551e-01, 1.34448847e-04],
       [9.99997881e-01, 2.11928102e-06],
       [9.97914068e-01, 2.08593214e-03],
       [9.99609179e-01, 3.90821247e-04],
       [2.23216260e-02, 9.77678374e-01],
       [1.47320254e-02, 9.85267975e-01],
       [1.16408255e-01, 8.83591745e-01],
       [9.99999197e-01, 8.02904032e-07]])
```

```
gnb.predict_proba(x_test)[0:10,1]
```

```
array([9.40462631e-01, 7.40639066e-06, 1.34448847e-04, 2.11928102e-06,  
       2.08593214e-03, 3.90821247e-04, 9.77678374e-01, 9.85267975e-01,  
       8.83591745e-01, 8.02904032e-07])
```

```
ypredict1 = gnb.predict_proba(x_test)[:,-1]
```

```
plt.hist(ypredict1,bins=6)  
plt.title('histogram')  
plt.xlabel('probabilities')  
plt.ylabel('frequency')
```



```
scores = cross_val_score(gnb,x_train,y_train,cv=6,scoring='accuracy')  
print('cross validation scores:{}'.format(scores))
```

```
cross validation scores:[0.79681728 0.79227053 0.80591077 0.8076158  0.79562251 0.78538943]
```

```
print('average cross validation scores:{}'.format(scores.mean()))
```

---

```
average cross validation scores:0.7972710529477408
```

---